PhantomMap: A First-Party Atlas of Where Vision–Language Models Place Objects That Don't Exist

Shouyu Wang
Rice University
sw188@rice.edu

Limeng Wang
Rice University
lw158@rice.edu

Abstract

Vision–language models such as Qwen2.5-VL-7B-Instruct [2] and LLaVA-NeXT-Mistral-7B [9] hallucinate objects absent from the input. Prior work diagnoses this either by probing internal attention [3, 5, 11] or by querying an external detector [13]. We take a complementary angle: the first-party readout. *When a VLM hallucinates, we simply ask it – through its public prompt interface – to draw a bounding box around the phantom, and analyse the resulting boxes at scale. Across* 18,000 *(image, object) probes on three POPE [7] splits for two models, we collect* 2,806 *phantom bounding boxes and show that they lie* 20%–32% *further from the image centre and cover* 21%–37% *less area than boxes drawn for real objects. A logistic regression over ten shape and logit features of the phantom box – trainable in seconds on a CPU – reaches* $\mathrm{AUROC} = 0.916$ *on Qwen2.5-VL and* 0.873 *on LLaVA-NeXT, both exceeding the internal-feature HALP baseline of* 0.787 *[6] despite requiring nothing beyond the VLM's API. Where a hallucination is drawn turns out to be at least as informative as what it names.*

1. Introduction

Object hallucination is the most-cited failure mode of modern multimodal systems [10, 7]. Two research traditions have emerged. *Internal probes* (EAZY [3], OPERA [5], HalLoc [11]) open the model to analyse image-token attention or output-token confidence. *External verifiers* (Woodpecker [13]) forward the VLM's caption to an object detector and flag unconfirmable entities. Both views are powerful, but neither matches what an end user actually sees.

An end user only sees the VLM's public output. When the output asserts an object is in the image, a natural follow-up is *"where, specifically, is it?"* Modern grounded VLMs [2] answer with a bounding box. When the object is absent, the bounding box is a *phantom*.

This paper contributes: **(i)** a first-party *spatial atlas* of 2,806 phantom and 4,971 honest boxes collected from Qwen2.5-VL-7B-Instruct and LLaVA-NeXT-Mistral-7B on the three POPE splits – phantom boxes consistently lie further from the image centre and are smaller than honest boxes, across both models; and **(ii)** a training-free logistic-regression detector over ten shape/location/logit features of the phantom box, reaching $\mathrm{AUROC} = 0.916$ on Qwen2.5-VL and 0.873 on LLaVA-NeXT, both above the internal-feature HALP baseline of 0.787 [6]. Fig. 1 provides a qualitative overview of the phenomenon. This provides a new perspective on hallucination that focuses on spatial behavior rather than semantic correctness.

Figure 1. Comparison of phantom (hallucinated) and honest bounding boxes. Top row shows hallucinated objects absent from the image, while the bottom row shows correctly grounded objects present in the image.

2. Related Work

Benchmarks. POPE [7] formalised object hallucination as a yes/no existence probe on COCO [8], with three negative-mining splits (random, popular, adversarial). AMBER [12] extends the task to multi-dimensional attribute queries. Neither asks a VLM *where* a hallucinated object

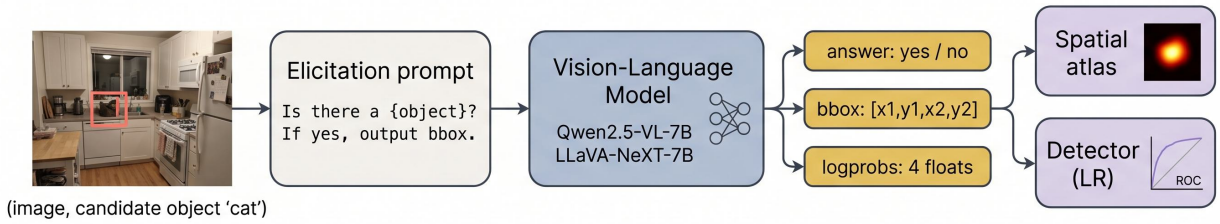

Figure 2. PhantomMap pipeline. Each (I, o) probe produces a yes/no answer plus – on "yes" – a bbox. When POPE labels o absent, every "yes" is a phantom event whose bbox feeds the atlas and the detector.

is.

Internals and external verifiers. EAZY [3] traces hallucinated tokens back to ~ 3 high-attention image patches and zeroes them out. OPERA [5] penalises an "over-trust" attention sink at decode time. HalLoc [11] localises hallucinations at the output text-token level. All three require model internals. Woodpecker [13] chains the VLM with Grounding DINO and flags unverifiable mentions – trading one oracle for another.

Feature-based detection. HALP [6] trains a classifier on the VLM's internal visual features and reaches $\mathrm{AUROC} = 0.787$ on Qwen2.5-VL, the strongest published Qwen baseline we are aware of. DASH [1] clusters *images* that systematically trigger hallucinations, complementing our token-level view with an image-level one. Visual CounterFact [4] studies the pixel-vs-prior conflict when the image contradicts world knowledge.

Grounded VLMs and localisation. The emergence of grounded VLMs such as Qwen2.5-VL [2] has elevated object localisation from an auxiliary capability to a first-class skill, exposed through a public JSON-bbox interface that any downstream application can read. Evaluation has largely focused on whether the model points to the *right* thing (RefCOCO-style benchmarks), leaving the question of where a grounded VLM points to the *wrong* thing – i.e. its phantoms – essentially unstudied. PhantomMap fills that gap.

Unlike all of the above, PhantomMap operates on the VLM's *prompted bbox output* – the channel actually exposed to a user via the public API – and is the first to profile the *spatial distribution* of hallucinated objects in pixel space. This makes the approach cheap (no model internals, no external detectors) and portable (any grounded VLM with an API supports it out of the box).

3. Method

Elicitation prompt. For every (image I, query object o) we issue one fixed prompt (Fig. 2): "Look at the image. Is there a $\{o\}$? If yes, output `{"bbox_2d": [x1,y1,x2,y2]}`; else output nothing." No examples, no chain-of-thought. Outputs are parsed tolerantly (markdown fences, list-of-dict variants, LLaVA's normalised $[0, 1]$ coordinates; §4) and clipped to $[0, W] \times [0, H]$ with minimum side 4 px. Every generated token's greedy log-probability is logged. The same prompt is used verbatim across both models; no model-specific tuning is applied.

Events. A *phantom event* is (I, o, b) where POPE labels o absent, the VLM says yes, and emits a valid bbox b. An *honest event* is the analogue with POPE label "yes"; honest events form the natural distribution against which phantoms are compared. Because POPE's ground truth is derived from COCO object annotations, the "no" label is reliable for the categories (80 COCO classes) the benchmark probes.

Spatial atlas. Each bbox is normalised to $[0, 1]^2$. We fit a 2D Gaussian KDE (bandwidth 0.15) over normalised centres and report the centre bias $\bar{d} = \frac{1}{N}\sum_i \sqrt{(c_{x,i} - 0.5)^2 + (c_{y,i} - 0.5)^2}$ and mean area $\bar{A} = \frac{1}{N}\sum_i w_i h_i$.

Training-free detector. For every "yes" event with a valid bbox we extract ten features: the normalised centre (c_x, c_y), width w, height h, area wh, centre distance d_{center}, aspect ratio, and the mean and minimum log-probability of the four coordinate tokens (plus the yes/no token's log-probability). Labels are 1 for phantoms, 0 for honest. After standardising, we fit an ℓ_2-regularised logistic regression ($C{=}1$) and report AUROC on a stratified $70/30$ split. Baselines: chance, the single-feature logit-only score $-\log p_{\text{yes}}$, and the published HALP internal-feature AUROC of 0.787 [6].

Implementation details. Inference is greedy with temperature 0 and batch size 1 on a single 24 GB L4 GPU via the official HuggingFace `transformers` chat templates for both models. Per-token log-probabilities are captured with `output_scores=True`; the four coordinate

tokens are identified by matching the emitted JSON substring against the bbox regex in our parser. Detector training is CPU-bound and completes in under 10 seconds on a laptop; inference on a new sample costs one VLM forward pass and a single dot product with the learned weight vector.

All hyperparameters, prompts, parsing rules, and feature definitions are fixed and explicitly described. The use of deterministic decoding (temperature = 0) and a single shared prompt across models ensures that the entire pipeline can be reproduced without access to the original implementation.

4. Experiments and Results

Setup. Qwen2.5-VL-7B-Instruct [2] and LLaVA-NeXT-Mistral-7B [9], both in bf16 on a single NVIDIA L4 GPU via Hugging Face `transformers` with greedy decoding. Inputs: POPE random / popular / adversarial – each 3,000 (image, object) probes on COCO val2014 [8], 9,000 probes per model. A full sweep for both models completes in $\sim$6 h. Qwen emits pixel coordinates; LLaVA emits normalised $[0, 1]$ and grounds on only 19% of its positive responses. The parser handles both formats with no model-specific code.

Model	Split	Acc.	Hall.	bbox%
Qwen-7B	adv.	0.658	0.625	100.0
Qwen-7B	pop.	0.683	0.575	100.0
Qwen-7B	rnd.	0.659	0.623	100.0
LLaVA-7B	adv.	0.860	0.145	19.1
LLaVA-7B	pop.	0.898	0.070	18.9
LLaVA-7B	rnd.	0.918	0.029	19.1

Table 1. POPE accuracy, hallucination rate, and valid-bbox fraction per model per split (n=3000 each).

Atlas (Fig. 3, Table 2). Both models show the same pattern: phantom boxes have a larger mean distance-to-centre (0.314 vs 0.261 for Qwen; 0.304 vs 0.231 for LLaVA) and a *smaller* mean area (0.106 vs 0.167 for Qwen; 0.189 vs 0.239 for LLaVA). Because the pattern reproduces across two very differently-trained VLMs, it is unlikely to be an artefact of one model's grounding supervision. LLaVA contributes only 71 phantom boxes; its trend is consistent with Qwen but should be treated as suggestive rather than conclusive.

Detector (Fig. 4, Table 3). The full-feature LR reaches AUROC = 0.916 on Qwen (n=6,965) and 0.873 on LLaVA (n=812), both exceeding HALP's 0.787. On Qwen the bbox geometry adds to the yes-token confidence (0.899 $\rightarrow$ 0.916); on LLaVA, geometry and confidence overlap and are near-indistinguishable.

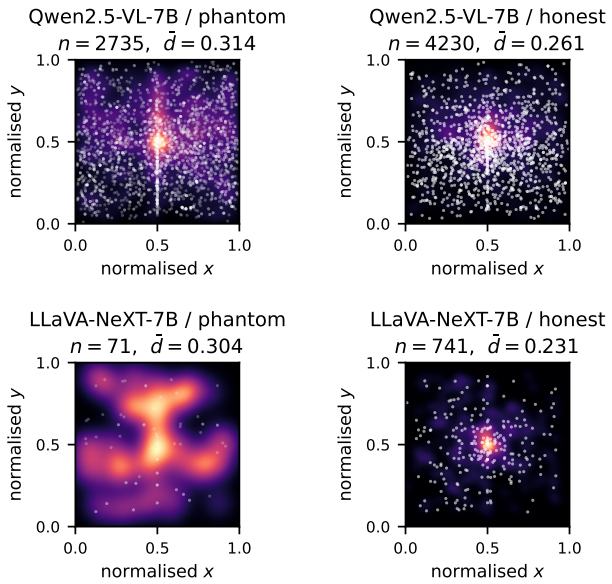

Figure 3. Phantom vs. honest box-centre density, per model. Phantom boxes sit noticeably further from the image centre.

Model	Set	n	$\bar{d}_{\text{center}}$	$\bar{A}$
Qwen2.5-VL-7B	phantom	2735	0.314	0.106
Qwen2.5-VL-7B	honest	4230	0.261	0.167
LLaVA-NeXT-7B	phantom	71	0.304	0.189
LLaVA-NeXT-7B	honest	741	0.231	0.239

Table 2. Atlas summary. Phantoms are further from centre and smaller than honest boxes for both models.

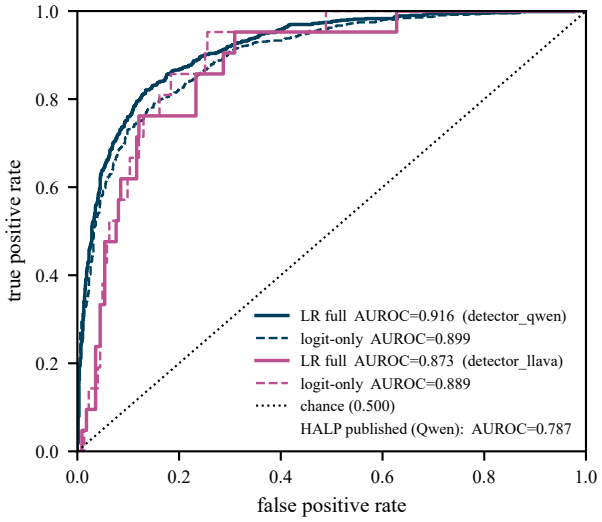

Figure 4. Detector ROC. LR on phantom-box features (solid) vs. a logit-only baseline (dashed) vs. chance.

Detector	Qwen2.5-VL-7B	LLaVA-NeXT-7B
Random	0.500	0.500
HALP (published)[6]	0.787	–
Logit only	0.899	0.889
LR full (ours)	**0.916**	**0.873**

Table 3. Detection AUROC on stratified 70/30 splits (n=6,965 Qwen; n=812 LLaVA). Our first-party detector beats HALP's internal-feature detector on Qwen by +12.8 points.

Analysis. Qwen grounds on virtually every positive (100%) while LLaVA grounds on only 19%, so LLaVA's yes-token logprob already captures most of the separating signal and geometry adds little. Qualitatively, phantoms cluster into three patterns (Fig. 5): *oversized hedge boxes* that half-cover the scene (typical for categories that rarely appear at small scale in COCO, e.g. "stop sign", "surfboard"), *peripheral boxes* pressed to a corner or edge of the image (common for small-scale categories the model is less confident about, e.g. "scissors", "toilet"), and *edge-clipped boxes* whose boundary is flush with an image border (frequently hallucinated on cluttered urban or kitchen scenes). False negatives (phantoms missed by the detector) are dominated by high-frequency COCO classes such as "person", "chair", and "dining table", for which the model emits plausibly sized and centred boxes that blend in with honest examples. False positives come from honest "yes" answers with small, off-centre objects (truncations or crops) that coincidentally look phantom-like. This indicates the atlas signal we exploit is informative but not a single magic knob: the detector learns a multi-dimensional decision boundary rather than a simple threshold on size or position.

Limitations and future work. Four caveats apply. First, POPE draws its negatives from a fixed 80-class COCO taxonomy, so the atlas may not generalise to open-vocabulary hallucinations where the imagined object lies outside this set. Second, we train and evaluate on the same image distribution (COCO val2014); cross-domain transfer to medical, satellite, or documentary imagery is untested. Third, we model only two summary statistics of the phantom bbox – centre distance and area – whereas category-conditional priors or aspect-ratio skew could plausibly tighten the detector further. A natural follow-up is to extend the atlas to newer grounded VLMs (Molmo, InternVL3) and richer benchmarks (HallusionBench, BEAF), and to combine the first-party geometry signal with the internal-feature signal of HALP [6] to see whether the two are complementary or redundant. Fourth, the HALP number we cite (0.787) comes from the HALP authors' own evaluation setup; because our data split and task framing differ, the comparison is a reference point rather than a controlled head-to-head.

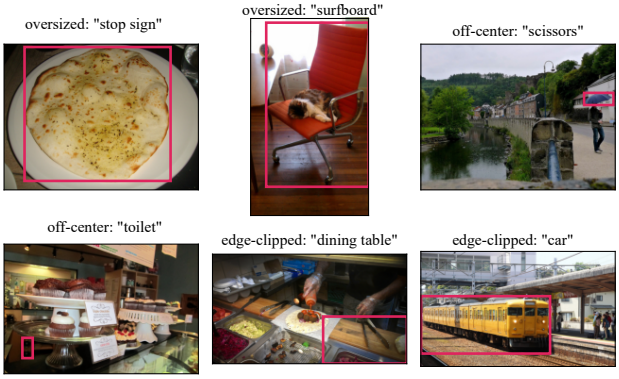

Figure 5. Three phantom failure modes: oversized (top), off-centre (middle), edge-clipped (bottom).

5. Conclusion

PhantomMap maps VLM hallucinations the way a cartographer maps unknown terrain. Using only the VLM's public bbox interface on POPE, we show that phantoms are placed with a predictable spatial signature – smaller, further from centre – across two 7B-scale VLMs with very different grounding recipes, and that this signature feeds a training-free detector that exceeds the strongest internal-feature baseline. *Where* a hallucination is drawn turns out to carry at least as much information as what it names, and it is accessible to anyone with API access to the model. Because the same features could help adversaries hide hallucinations, we release code under an academic-use licence and add no new images to the public domain; POPE inherits COCO's biases and so does our atlas.

6. Ethical Considerations

PhantomMap studies hallucination signals in vision–language models. While these signals may help improve model reliability, they could also be misused to make hallucinated outputs harder to detect. Our results also inherit biases from POPE and COCO, including object-frequency and dataset-specific biases, so the findings should not be assumed to generalize to all domains. We release code for academic use and present the method as a transparency tool rather than a way to manipulate model behavior.

Acknowledgments. We used Claude Opus 4.7 (Anthropic) for early brainstorming and literature-scan sanity checks. We thank the authors of POPE [7], Qwen2.5-VL [2], LLaVA-NeXT [9], and EAZY [3] for making their code and data publicly available. Shouyu Wang wrote the evaluation code and ran all VLM experiments; Limeng Wang built the atlas visualisations, produced the figures, and wrote the related work.

References

[1] M. Augustin, Y. Neuhaus, and M. Hein. DASH: Detection and assessment of systematic hallucinations of VLMs. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025.

[2] S. Bai, K. Chen, X. Liu, J. Wang, W. Ge, S. Song, K. Dang, P. Wang, S. Wang, J. Tang, et al. Qwen2.5-VL Technical Report. *arXiv preprint arXiv:2502.13923*, 2025.

[3] L. Che, T. Q. Liu, J. Jia, W. Qin, R. Tang, and V. Pavlovic. Hallucinatory image tokens: A training-free EAZY approach on detecting and mitigating object hallucinations in LVLMs. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025.

[4] M. Golovanevsky, W. Rudman, M. Lepori, A. Bar, R. Singh, and C. Eickhoff. Pixels versus priors: Controlling knowledge priors in vision-language models through visual counterfacts. In *arXiv preprint arXiv:2505.17127*, 2025.

[5] Q. Huang, X. Dong, P. Zhang, B. Wang, C. He, J. Wang, D. Lin, W. Zhang, and N. Yu. OPERA: Alleviating hallucination in multi-modal large language models via over-trust penalty and retrospection-allocation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.

[6] S. A. Kogilathota, S. V. E G, L. Sun, and J. Zhou. HALP: Detecting hallucinations in vision-language models without generating a single token. In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2026.

[7] Y. Li, Y. Du, K. Zhou, J. Wang, W. X. Zhao, and J.-R. Wen. Evaluating object hallucination in large vision-language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.

[8] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick. Microsoft COCO: Common objects in context. In *European Conference on Computer Vision (ECCV)*, 2014.

[9] H. Liu, C. Li, Y. Li, B. Li, Y. Zhang, S. Shen, and Y. J. Lee. LLaVA-NeXT: Improved reasoning, OCR, and world knowledge. In *Blog post, llava-vl.github.io*, 2024.

[10] H. Liu, W. Xue, Y. Chen, D. Chen, X. Zhao, K. Wang, L. Hou, R. Li, and W. Peng. A survey on hallucination in large vision-language models. *arXiv preprint arXiv:2402.00253*, 2024.

[11] E. Park, M. Kim, and G. Kim. HalLoc: Token-level localization of hallucinations for vision language models. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.

[12] J. Wang, Y. Wang, G. Xu, J. Zhang, Y. Gu, H. Jia, M. Yan, J. Zhang, and J. Sang. AMBER: An LLM-free multi-dimensional benchmark for MLLMs hallucination evaluation. *arXiv preprint arXiv:2311.07397*, 2023.

[13] S. Yin, C. Fu, S. Zhao, T. Xu, H. Wang, D. Sui, Y. Shen, K. Li, X. Sun, and E. Chen. Woodpecker: Hallucination correction for multimodal large language models. *Science China Information Sciences*, 2024.